

# Lab Setup

SSH or Remote Desktop clients:

- SSH Access Software (recommended): PuTTY for Windows can be downloaded from [www.putty.org](http://www.putty.org)
  - Alternatively you may use a provided browser-based SSH option
- Remote Desktop Software: Download NoMachine now for best performance from [www.nomachine.com/download](http://www.nomachine.com/download)
  - Alternatively you may use a VNC client or the provided browser-based VNC option

Connection Instructions :

1. Navigate to [nvlabs.qwiklab.com](http://nvlabs.qwiklab.com);
2. Login or create a new account;
3. Select the *Instructor-Led Hands-on Labs Class*;
4. Find the lab called *Applied Deep Learning for Vision and Natural Language with Torch7*;
  - select it, click Select, and finally click Start
5. After a short wait, lab instance Connection information will be shown

Please ask Lab Assistants for help!

---

class: center, middle

## Torch 7: Applied Deep Learning for Vision and Natural Language

Nicholas Leonard

<https://github.com/nicholas-leonard>

Element Inc.

April 6, 2016

---

## Agenda

1. Introduction
2. Packages
3. Tensors
4. Logistic Regression (Exercise 1)

5. Deep Learning
  6. Multi-Layer Perceptron (Exercise 2)
  7. Convolutional Neural Network (Exercise 3)
  8. Recurrent Neural Network (Exercise 4)
- 

## Introduction

My background:

- 2003-2008 : Bac. Degree in Comp. Sci. at Royal Military College of Canada ;
  - 2008-2012 : Army Signals Officer :
    - management (office politics, emails), no code, no science ;
    - hobby : neural networks using Python and numpy ;
  - 2012-2014 : Master's Degree in Deep Learning at University of Montreal ;
    - LISA/MILA lab ;
    - Yoshua Bengio and Aaron Courville as co-directors ;
    - 2012-2013 : Python, Theano and Pylearn2 ;
    - 2013-today : Lua, Torch7 ;
  - 2014-today : research engineer at Element Inc. :
    - biometrics startup ;
    - deep learning on smart phones (Android, iOS) ;
    - open source contributions (Torch7).
- 

## Introduction - Lua

Why take the time to learn Lua :

- easy interface between low-level C/CUDA/C++ and high-level Lua ;
- light-weight : used for embedded systems ;
- tables :
  - can be used as lists, dictionaries, packages, classes and objects ;
  - make it easy to extend existing classes (at any level) ;
- fast for-loops (LuaJIT) ;
- closures ;

Example :

```
a = {1,2,a=3, print=function(self) print(self) end}
a:print() -- i.e. a.print(a)
```

Output :

```
{
  1 : 1
  2 : 2
  print : function: 0x417f11e0
  a : 3
}
```

---

## Introduction - Torch 7

What's up with Torch 7?

- a powerful N-dimensional array ;
- lots of routines for indexing, slicing, transposing, ... ;
- amazing interface to C, via LuaJIT ;
- linear algebra routines ;
- easy modular neural networks ;
- numeric optimization routines ;
- fast and efficient GPU support ;
- ports to iOS, Android and FPGA backends ;
- under development since October 2002 ;
- used by Facebook, Google [DeepMind], Twitter, NYU, ... ;
- documentation, tutorials, demos, examples ;

---

## Introduction - Useful Links

Main : <http://torch.ch/>

Cheatsheet: <https://github.com/torch/torch7/wiki/Cheatsheet>

Github: <https://github.com/torch/torch7>

Google Group for new users and installation queries:

<https://groups.google.com/forum/embed/?place=forum%2Ftorch7#!forum/torch7>

Advanced only: <https://gitter.im/torch/torch7>

---

## Packages

The Torch 7 distribution is made up of different packages, each its own github repository :

- **torch7/cutorch** : tensors, BLAS, file I/O (serialization), OOP, unit testing and cmd-line argument parsing ;
- **nn/cunn** : easy and modular way to build and train neural networks using `modules` and `criteria` ;
- **nngraph** : nn with support for more complicated graphs ;
- **optim** : optimization package for nn. Provides training algorithms like SGD, LBFGS, etc. Uses closures ;
- **trepl** : torch read-eval-print loop, Lua interpreter, `th>` ;
- **paths** : file system manipulation package ;
- **image** : for saving, loading, constructing, transforming and displaying images ;

Refer to the [torch.ch](http://torch.ch) website for a more complete list of official packages.

---

## Packages - Unofficial

Many more unofficial packages out there :

- **dpnn** : extensions to the nn library. ;
  - **rnn** : recurrent neural network library. Implements RNN, GRU, LSTM, BRNN, and RAM ;
  - **nnx/cunnx** : experimental neural network modules and criteria : `SpatialReSampling`, `SoftMaxTree`, etc. ;
  - **dataloader** : library for loading and iterating through datasets ;
  - **moses** : utility-belt library for functional programming in Lua, mostly for tables ;
  - **threads/parallel** : libraries for multi-threading or multi-processing ;
- 

## Tensors

Tensors are the main class of objects used in Torch 7 :

- An N-dimensional array that views an underlying `Storage`;
- Different Tensors can share the same `Storage`;
- Different types : `FloatTensor`, `DoubleTensor`, `IntTensor`, `CudaTensor`, ...;
- Implements many Basic Linear Algebra Sub-routines (BLAS) :
  - `torch.addmm` : matrix-matrix multiplication ;
  - `torch.addmv` : matrix-vector multiplication ;
  - `torch.addr` : outer-product between vectors ;
  - etc.
- Supports random initialization, indexing, transposition, sub-tensor extractions, and more ;
- Most operations for Float/Double are also implemented for Cuda Tensors (via **cutorch**) ;

---

## Tensors - Initialization

A 2x3 Tensor :

```
th> a = torch.FloatTensor(2,3)
-- initialized with garbage content (whatever was already there)
th> a
  8.6342e+19  4.5694e-41  8.6342e+19
  4.5694e-41  0.0000e+00  0.0000e+00
[torch.FloatTensor of size 2x3]
```

Fill with ones :


```
th> a:fill(1)
  1  1  1
  1  1  1
[torch.FloatTensor of size 2x3]
```

Random uniform initialization :

```
th> a:uniform(0,1) -- random uniform between 0 and 1
  0.6323  0.9232  0.2930
  0.8412  0.5131  0.9101
[torch.FloatTensor of size 2x3]
```

---

## Tensors - BLAS



.center[ ]

Tensors are all about basic linear algebra. Let's multiply an `input` and a `weight` matrix into an `output matrix` :

```
batchSize, inputSize, outputSize = 4, 2, 3
```

```
input = torch.FloatTensor(batchSize, inputSize):uniform(0,1)
weight = torch.FloatTensor(outputSize, inputSize):uniform(0,1)
output = torch.FloatTensor(batchSize, outputSize):zero()
-- matrix matrix multiply :
output:addmm(0, output, 1, input, weight:t())
```

This is a common operation used by the popular `nn.Linear` module.

---

## Tensors - CUDA

Previous matrix-matrix multiply using CUDA :

```
require 'cutorch'
input = input:cuda() or torch.CudaTensor(batchSize, inputSize):uniform(0,1)
weight = weight:cuda() or torch.CudaTensor(outputSize, inputSize):uniform(0,1)
output = output:cuda() or torch.CudaTensor(batchSize, outputSize):zero()
-- matrix matrix multiply :
output:addmm(0, output, 1, input, weight:t())
```

So basically, no difference except for use of `torch.CudaTensor`.

---

## Neural Network library

The `nn` package :

- implements feed-forward neural networks ;
- neural networks form a computational flow-graph of transformations ;
- backpropagation is gradient descent using the chain rule ;

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$

Two abstract classes :

- `nn.Module` : differentiable transformations of input to output ;
- `nn.Criterion` : cost function to minimize. Outputs a scalar loss;

Let's use it to build a simple logistic regressor...

---

## Logistic Regression - Module

A binary logistic regressor `Module` with 2 input units and 1 output.


```
require 'nn'
module = nn.Sequential()
module:add(nn.Linear(2, 1))
module:add(nn.Sigmoid())
```

The above implements :

$$y = \sigma(Wx + b)$$

where the sigmoid (logistic function) is defined as :

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



---

## Logistic Regression - Criterion and Data

A binary cross-entropy `Criterion` (which expects 0 or 1 valued targets) :

```
criterion = nn.BCECriterion()
```

The BCE loss is defined as :

$$-\sum_i [t_i \log(y_i) + (1 - t_i) \log(1 - y_i)]$$

Some random dummy dataset with 10 samples:

```
inputs = torch.Tensor(10,2):uniform(-1,1)
```

```
targets = torch.Tensor(10):random(0,1)
```

---

## Logistic Regression - Training

Function for one epoch of stochastic gradient descent (SGD)

```
require 'dpnn'
function trainEpoch(module, criterion, inputs, targets)
  for i=1,inputs:size(1) do
    local idx = math.random(1,inputs:size(1))
    local input, target = inputs[idx], targets:narrow(1,idx,1)
    -- forward
    local output = module:forward(input)
    local loss = criterion:forward(output, target)
    -- backward
    local gradOutput = criterion:backward(output, target)
    module:zeroGradParameters()
    local gradInput = module:backward(input, gradOutput)
    -- update
    module:updateGradParameters(0.9) -- momentum (dpnn)
    module:updateParameters(0.1) --  $W = W - 0.1 * dL/dW$ 
  end
end
```

---

## Exercise 1 : Logistic Regression

Modify the `logistic-regression.lua` script to do the following :

1. Train for 100 epochs;
2. Each call to `trainEpoch` prints the mean error of that epoch;
3. Bonus : Use `Softmax + ClassNLLCriterion` instead of `Sigmoid + BCECriterion`.

Should output something like :

```
$ th logistic-regression.lua
...
Epoch 99 : mean loss = 0.038046
Epoch 100 : mean loss = 0.036116
```

Time : 10 min.

Solution found in `solution/logistic-regression.lua`.

---



## Exercise 1 : Take-away points

Like Python, Lua is relatively easy to use;

Stochastic Gradient Descent is simple: \* forward, backward, update; \* error is minimized after multiple iterations through data;

Can use BCE or NLL to minimize error of binary classification problems;

Logistic regression models are very simple models.

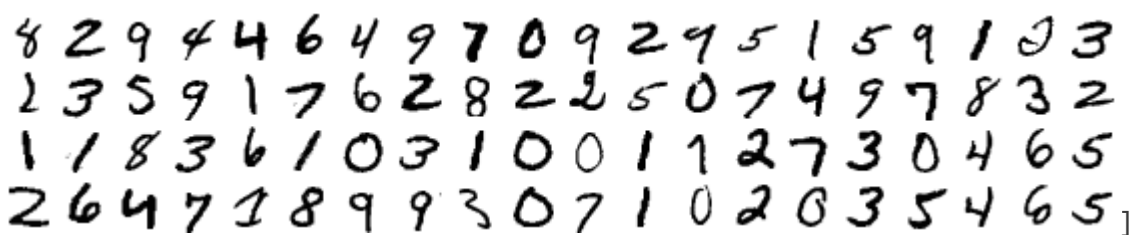
---

## Deep Learning

What is deep learning?

- **collection of techniques** to improve the optimization and **generalization** of neural networks :
    - rectified linear units ;
    - dropout ;
    - batch normalization ;
    - weight decay regularization ;
    - momentum learning ;
  - **stacking layers** of transformations to create successively more abstract **levels of representations** ;
    - depth over breadth ;
    - deep multi-layer perceptrons ;
  - **shared parameters** :
    - convolutional neural networks ;
    - recurrent neural networks ;
  - **technological improvements** :
    - massively parallel processing : GPUs, CUDA ;
    - fast libraries : torch, cudnn, cuda-convnet, theano, tensorflow;
- 

## Deep Learning - MNIST dataset

  
The image displays a 4x15 grid of handwritten digits from the MNIST dataset. The digits are arranged in four rows and fifteen columns. The first row contains: 8, 2, 9, 4, 4, 6, 4, 9, 7, 0, 9, 2, 9, 5, 1, 5, 9, 1, 0, 3. The second row contains: 1, 3, 5, 9, 1, 7, 6, 2, 8, 2, 2, 5, 0, 7, 4, 9, 7, 8, 3, 2. The third row contains: 1, 1, 8, 3, 6, 1, 0, 3, 1, 0, 0, 1, 1, 2, 7, 3, 0, 4, 6, 5. The fourth row contains: 2, 6, 4, 7, 1, 8, 9, 9, 3, 0, 7, 1, 0, 2, 0, 3, 5, 4, 6, 5. The digits are centered horizontally within each row.

`dataload` package makes it easy to obtain the MNIST dataset :

```
local dl = require 'dataloader'
local trainset, validset, testset = dl.loadMNIST()
```

Each *batch* is a 4D `inputs` tensor and a `targets` 1D tensor:

```
inputs, targets = trainset:sample(32)
print(inputs:size(), targets:size())
32 1 28 28
[torch.LongStorage of size 4]


32
[torch.LongStorage of size 1]
```

background-image:

[url\(https://raw.githubusercontent.com/nicholas-leonard/slides/master/we-need-to-go-deeper.jpg\)](https://raw.githubusercontent.com/nicholas-leonard/slides/master/we-need-to-go-deeper.jpg)

## Deep Learning

### Multi-Layer Perceptron



An MLP is a stack of non-linear layers :

- each layer is an affine transform (`Linear`) followed by a transfer function (`Tanh`, `ReLU`, `SoftMax`) ;
- parameters (`weight`, `bias`) are found in the `Linear` module ;
- parameters are varied to fit the data ;
- transfer functions help to model complex relationships between input and output (non-linear);

---

## Multi-Layer Perceptron - Module and Criterion

An MLP with 2 layers of hidden units :

```
module = nn.Sequential()
module:add(nn.Convert()) -- casts input to model type (float -> double)
module:add(nn.Collapse(3)) -- collapse 3D to 1D
module:add(nn.Linear(1*28*28, 200))
module:add(nn.Tanh())
module:add(nn.Linear(200, 200))
module:add(nn.Tanh())
module:add(nn.Linear(200, 10))
module:add(nn.LogSoftMax()) -- for classification problems
```

Negative Log-Likelihood (NLL) Criterion :

```
criterion = nn.ClassNLLCriterion()
```

---

## Multi-Layer Perceptron - Cross-validation

A function to evaluate performance on the validation set :

```
require 'optim'
cm = optim.ConfusionMatrix(10)
function classEval(module, inputs, targets)
    cm:zero()
    for idx=1,inputs:size(1) do
        local input, target = inputs[idx], targets:narrow(1,idx,1)
        local output = module:forward(input)
        cm:add(output, target)
    end
    cm:updateValid()
    return cm.totalValid
end
```

Measure model's ability to *generalize* to new data.

---

## Multi-Layer Perceptron - Early-Stopping

Early-stopping on the validation set :

```

bestAccuracy, bestEpoch = 0, 0
wait = 0
for epoch=1,300 do
    trainEpoch(module, criterion, trainInputs, trainTargets)
    local validAccuracy = classEval(module, validInputs, validTargets)
    if validAccuracy > bestAccuracy then
        bestAccuracy, bestEpoch = validAccuracy, epoch
        torch.save("/path/to/saved/model.t7", module)
        print(string.format("New maxima : %f @ %f", bestAccuracy, bestEpoch))
        wait = 0
    else
        wait = wait + 1
        if wait > 30 then break end
    end
end
end

```

Early-stops when no new maxima has been found for 30 consecutive epochs.

---

## Exercise 2 : Multi-Layer Perceptron

Modify the `multi-layer-perceptron.lua` script to do the following :

1. take options from the command-line;
2. add `Dropout` between hidden layers;
3. Bonus : write script to evaluate saved model on test set;
4. Bonus : make the number of hidden layers a hyper-parameter.

Time : 10 min.

Evaluation script should output something like :

```

th evaluate-mlp.lua
Test accuracy=0.977000

```

Solution found in `solution/[train|evaluate]-mlp.lua`.

---

## Exercise 2 : Take-away points

Easier to try different hyper-parameters (options) from the cmd-line;


Dropout can help with generalization, but increases convergence time;

Experimentation usually requires two scripts and 3 datasets : \* training : optimize model on *training set*, early-stop on *validation set*; \* evaluation : measure performance on *test set*;

Deeper models (more layers) can help with generalization.

# Convolutional Neural Network

.center[




]

CNNs are often stacks of meta-layers each made from 3 layers :

1. **convolution** : convolve a kernel over the image along height and width axes ;
2. **transfer function** : a non-linearity like  $\tanh$  or  $\text{ReLU}$  ;
3. **sub-sampling** : reduce the size (height x width) of feature maps by pooling them spatially;

## Convolutional Neural Network - Convolution



.center[

]

Convolution modules typically have the following arguments :

- `padSize` : how much zero-padding to add around the input image ;
- `inputSize` : number of input channels (e.g. 3 for RGB image) ;
- `outputSize` : number of output channels (number of filters) ;
- `kernelSize` : height and width of the kernel ;
- `kernelStride` : step-size of the kernel (typically 1) ;

Parameters of the convolution (i.e. the kernel) :

- `weight` : 4D Tensor `outputSize x inputSize x kernelSize x kernelSize` ;
  - `bias` : 1D Tensor `outputSize` ;
- 

## Convolutional Neural Network - SpatialConvolution


`SpatialConvolution` with 3 input and 4 output channels using a 5x5 kernel on a 12x12 image :

```
input = torch.rand(3,12,12)
conv = nn.SpatialConvolution(3,4,5,5)
output = conv:forward(input) -- size is 4 x 8 x 8
```

Now with 2 pixels of padding on each side:

```
conv = nn.SpatialConvolution(3,4,5,5,1,1,2,2)
output = conv:forward(input) -- size is 4 x 12 x 12
```

Learns filters like :



`.center[`

---

## Convolutional Neural Network - Sub-sampling

Sub-sampling modules :

- typically max-pooling is used : `SpatialMaxPooling` ;
- makes the model more invariant to translation ;
- reduces the size of the spatial dimensions ;

`SpatialMaxPooling` to pool inputs in a 2x2 area with a stride of 2:

```
input = torch.range(1,16):double():resize(1,4,4)
```

```
pool = nn.SpatialMaxPooling(2,2,2,2)
output = pool:forward(input)
print(input, output)
(1,...) =
  1   2   3   4
  5   6   7   8
  9  10  11  12
 13  14  15  16
[torch.DoubleTensor of size 1x4x4]

(1,...) =
  6   8
 14  16
[torch.DoubleTensor of size 1x2x2]
```

---

## Convolutional Neural Network - MNIST

Convolutional Neural Network for the MNIST dataset :

```
cnn = nn.Sequential()
-- 2 conv layers :
cnn:add(nn.Convert())
cnn:add(nn.SpatialConvolution(1, 16, 5, 5, 1, 1, 2, 2))
cnn:add(nn.ReLU())
cnn:add(nn.SpatialMaxPooling(2, 2, 2, 2))
cnn:add(nn.SpatialConvolution(16, 32, 5, 5, 1, 1, 2, 2))
cnn:add(nn.ReLU())
cnn:add(nn.SpatialMaxPooling(2, 2, 2, 2))
-- 1 dense hidden layer :
outsize = cnn:outsize{1,1,28, 28} -- output size of convolutions
cnn:add(nn.Collapse(3))
cnn:add(nn.Linear(outsize[2]*outsize[3]*outsize[4], 200))
cnn:add(nn.ReLU())
-- output layer
cnn:add(nn.Linear(200, 10))
cnn:add(nn.LogSoftMax())
```

---

## Convolutional Neural Network - Print Module

The `cnn` looks like this :

```
print(cnn)
nn.Sequential {
  [input -> (1) -> (2) -> (3) -> (4) -> (5) -> (6) -> (7) -> (8) -> (9) -> (10) ->
  (11) -> output]
  (1): nn.Convert
  (2): nn.SpatialConvolution(1 -> 16, 5x5, 1,1, 2,2)
  (3): nn.ReLU
```

```
(4): nn.SpatialMaxPooling(2,2,2,2)
(5): nn.SpatialConvolution(16 -> 32, 5x5, 1,1, 2,2)
(6): nn.ReLU
(7): nn.SpatialMaxPooling(2,2,2,2)
(8): nn.Linear(1568 -> 200)
(9): nn.ReLU
(10): nn.Linear(200 -> 10)
(11): nn.LogSoftMax
}
```

---

## Exercise 3 : Convolutional Neural Network

Use the `convolution-neural-network.lua` script.

Modify the script to do the following :

1. add `SpatialMaxPooling` to convolution layers;
2. add `[Spatial]BatchNormalization` to convolution and hidden layers ;
3. Bonus : efficiently save model to disk using `Serial`.

Try the `-cuda` flag.

Evaluate with `th solution/evaluate-mlp.lua -modelpath '/home/ubuntu/save/cnn.t7'`.

Time : 10 min.

Solution found in `solution/convolution-neural-network.lua`.

---

## Exercise 3 : Take-away points

CNNs have better performance than MLPs for images;

We can get a 4-10x speedup using NVIDIA GPUs;

`Serial` doesn't save `output` and `gradInput` to disk : 1.5 vs 9 MB;


Batch normalization can help with convergence and generalization;

Pooling reduces the width/height of representations.

---

## Recurrent Neural Network





.center[ s(t-1) ]

Simple RNN :

- for modeling sequential data like text, speech, videos ;
- 3 layers : input ( $v$ ), recurrent ( $u$ ) and output ( $w$ ) layer ;
- feed the previous state as input to next state ;
- long sequences suffer from exploding and vanishing gradient ;

## Recurrent Neural Network - Language Model


Maximize likelihood of next word given previous words (input -> target) :

1. we -> need
2. we, need -> to
3. we, need, to -> go
4. we, need, to, go -> deeper

Neural network language model (NNLM) :

- learn an embedding space of words ;
- each word is a vector of parameters ;
- embedding space is implemented using `LookupTable` ;
- embedding space is a weight matrix of size `vocabSize x embedSize` ;

## Recurrent Neural Network - BPTT



.center[ ]

Back-propagation through time :

- forward-propagate for  $T$  time-steps ;
- unfold network for  $T$  time-steps ;
- back-propagate through unfolded network ;
- accumulate parameter gradients (sum over time-steps) ;

## Recurrent Neural Network - Penn Tree Bank

Penn Tree Bank dataset :

- common benchmark for language models ;
- 10000 word vocabulary ;
- approx. 1 million words of text ;

Use **dataload** to get Penn Tree Bank dataset :

```
trainset, validset, testset = dl.loadPTB{batchsize,1,1}
```

Batch of 3 sample sequences of length 5:

```
print(inputs:t(), targets:t())
 238   808   951  1326  1477
 8311  7671  4749    49  2308
 3660  9800  4765  5375  6018
[torch.IntTensor of size 3x5]

  808   951  1326  1477  1692
 7671  4749    49  2308  9120
```

```
9800 4765 5375 6018 7523
[torch.IntTensor of size 3x5]
```

## Recurrent Neural Network - rnn

Use the `rnn` package to build an RNNLM.

A module that implements recurrence  $\{x[t], h[t-1]\} \rightarrow h[t]$ :

```
rm = nn.Sequential() -- input is {x[t], h[t-1]}
: add(nn.ParallelTable()
      : add(nn.LookupTable(10000, 200)) -- input layer (V)
      : add(nn.Linear(200, 200)) -- recurrent layer (U)
: add(nn.CAddTable())
: add(nn.Sigmoid()) -- output is h[t]
```

Wrap into a `Recurrence` module and add an output layer:

```
rnn = nn.Sequential()
: add(nn.Recurrence(rm, 200, 0))
: add(nn.Linear(200, 10000)) -- output layer (W)
: add(nn.LogSoftMax())
```

Wrap into a `Sequencer` to handle one sequence per `forward` call:

```
rnn = nn.Sequencer(rnn)
```

## Exercise 4 : Recurrent Language Model

Modify the `recurrent-language-model.lua` script to do the following :

1. use `LSTM`, `FastLSTM` or `GRU` instead of `Recurrence`;
2. train to reach 150 perplexity (PPL) on validation set;
3. use `evaluate-rnnlm.lua` to sample text from saved models.

```
nn.Serial @ nn.Sequential {
  [input -> (1) -> (2) -> (3) -> output]
  (1): nn.LookupTable
  (2): nn.SplitTable
  (3): nn.Sequencer @ nn.Recursor @ nn.Sequential {
    [input -> (1) -> (2) -> (3) -> (4) -> output]
    (1): nn.FastLSTM(200 -> 200)
    (2): nn.FastLSTM(200 -> 200)
    (3): nn.Linear(200 -> 10000)
```

```
(4): nn.LogSoftMax
}
}
```

Time : 15 min;

Solutions in `solution/train-rnnlm.lua`.

---

## Exercise 4 : Take-away points

LSTM and GRU can learn longer sequences (`seqlen`) than RNNs;

Stacking LSTM/GRU/RNNs can give even better results;

Optimizing hyper-parameters is a process of trial and error that takes time;

Using a learning rate schedule can help :

```
th recurrent-language-model.lua -progress -cuda -lstm -seqlen 20 -hiddensize
'{200,200}'
-batchsize 20 -startlr 1 -cutoff 5 -maxepoch 13
-schedule
'{[5]=0.5, [6]=0.25, [7]=0.125, [8]=0.0625, [9]=0.03125, [10]=0.015625, [11]=0.0078125, [1
2]=0.00390625}'
```

Evaluation loops through one continuous sequence;

---

## Exercise 4 : Generating text

Generating text using a model with 116 test PPL :

```
th evaluate-rnnlm.lua -xplogpath /home/ubuntu/save/rnnlm/rnnlm200x200.t7 -nsample
200
```

will result in something like this :

```
mr. jones contends that in his first popularity the organizations are
almost <unk> by their side <eos> it is particularly almost he adds <eos>
but it 's also a <unk> unsuccessfully <eos> i felt somebody wanted
westridge and financial commercial steelmakers <eos> an oct.
N bankruptcy-law article is comment <eos> <unk> <unk> <eos> there 's no
signs of internal mergers and companies involved in the life of light
insider business <eos> michael <unk> several years old and the mr.
buffett as a former abortions in the senate to abortion-rights its mind
```

```
is mr. straszheim 's remarks with a career replacement works <eos> <unk>  
education was named a director of this plastics media <eos>
```

Not bad for 1h of training!

---

## Recurrent Neural Network - Character LM

References : \* <https://github.com/karpathy/char-rnn> \* <https://github.com/hughperkins/char-lstm>

Text generated using char-level LM trained on 1M reddit comments:

```
<post>  
Diablo  
<comment score=1>  
I liked this game so much!! Hope telling that numbers' benefits and  
features never found out at that level is a total breeze  
because it's not even a developer/voice opening and rusher runs  
the game against so many people having noticeable purchases of selling  
the developers built or trying to run the patch to Jagex.  
</comment>
```

Looks good! Wait a second...

---

## Questions ?

.center[

